

2008 Special Issue

Training neural network classifiers for medical decision making: The effects of imbalanced datasets on classification performance[☆]

Maciej A. Mazurowski^{a,*}, Piotr A. Habas^a, Jacek M. Zurada^a, Joseph Y. Lo^b, Jay A. Baker^b,
Georgia D. Tourassi^b

^a Computational Intelligence Lab, Department of Electrical and Computer Engineering, University of Louisville, Louisville, KY 40292, USA

^b Duke Advanced Imaging Laboratories, Department of Radiology, Duke University Medical Center, Durham, NC 27705, USA

Received 9 August 2007; received in revised form 2 December 2007; accepted 11 December 2007

Abstract

This study investigates the effect of class imbalance in training data when developing neural network classifiers for computer-aided medical diagnosis. The investigation is performed in the presence of other characteristics that are typical among medical data, namely small training sample size, large number of features, and correlations between features. Two methods of neural network training are explored: classical backpropagation (BP) and particle swarm optimization (PSO) with clinically relevant training criteria. An experimental study is performed using simulated data and the conclusions are further validated on real clinical data for breast cancer diagnosis. The results show that classifier performance deteriorates with even modest class imbalance in the training data. Further, it is shown that BP is generally preferable over PSO for imbalanced training data especially with small data sample and large number of features. Finally, it is shown that there is no clear preference between oversampling and no compensation approach and some guidance is provided regarding a proper selection.

© 2007 Elsevier Ltd. All rights reserved.

Keywords: Classification; Feed-forward neural networks; Class imbalance; Computer-aided diagnosis

1. Introduction

In computer-aided decision (CAD) systems, computer algorithms are used to help a physician in diagnosing a patient. One of the most common tasks performed by a CAD system is the classification task where a label is assigned to a query case (i.e., a patient) based on a certain number of features (i.e., clinical findings). The label determines the query's membership in one of predefined classes representing possible diagnoses. CAD systems have been investigated and applied for the diagnosis of various diseases, especially for cancer. Some comprehensive reviews on the topic can be found in Kawamoto, Houlihan, Balas, and Lobach (2005), Lisboa and Taktak (2006) and Sampat, Markey, and Bovik (2005). CAD

systems rely on a wide range of classifiers, such as traditional statistical and Bayesian classifiers (Duda, Hart, & Stork, 2000), case-based reasoning classifiers (Aha, Kibler, & Albert, 1991), decision trees (Mitchell, 1997), and neural networks (Zhang, 2000). In particular, neural network classifiers are a very popular choice for medical decision making and they have been shown to be very effective in the clinical domain (Lisboa, 2002; Lisboa & Taktak, 2006).

To construct a classifier, a set of examples representing previous experience is essential. In general, the larger and more representative the set of available examples is, the better the classification of future query cases (Raudys & Jain, 1991). In the medical domain, however, there are several challenges and practical limitations associated with data collection. First, collecting data from patients is time consuming. Second, acquiring large volumes of data on patients representing certain diseases is often challenging due to the low prevalence of the disease. This is the case, for example, with CAD systems developed to support cancer screening. Cancer prevalence is particularly low among screening populations which results in class imbalance in the set of collected examples; a phenomenon

[☆] An abbreviated version of some portions of this article appeared in Mazurowski, Habas, Zurada, and Tourassi (2007) as part of the IJCNN 2007 Conference Proceedings, published under IEE copyright.

* Corresponding address: 407 Lutz Hall, University of Louisville, Louisville, KY 40222, USA. Tel.: +1 502 852 3165; fax: +1 502 852 3940.

E-mail address: maciej.mazurowski@louisville.edu (M.A. Mazurowski).

where one of the disease states is underrepresented. In addition, the clinical presentation of patients with the same disease varies dramatically. Due to this inherent variability, CAD systems are often asked to handle large numbers of features, many of which are correlated and/or of no significant diagnostic value. The issues described above (i.e., finite sample size, imbalanced datasets, and large numbers of potentially correlated features) can have a detrimental effect on the development and performance evaluation of typical CAD classifiers (Mazurowski et al., 2007).

Several investigators have addressed classification in the presence of these issues from both general machine learning and CAD perspectives. Most attention has been given to the effects of finite sample size (Beiden, Maloof, & Wagner, 2003; Chan, Sahiner, & Hadjiiski, 2004; Fukunaga & Hayes, 1989; Raudys, 1997; Raudys & Jain, 1991; Sahiner, Chan, Petrick, Wagner, & Hadjiiski, 2000; Wagner, Chan, Sahiner, & Petrick, 1997). The problem of large data dimensionality (i.e., large number of features) has been addressed in Hamamoto, Uchimura, and Tomita (1996) for neural networks and in Raudys (1997) and Raudys and Jain (1991) for other types of classifiers. In addition, researchers have examined the effect of finite sample on feature selection (Jain & Zongker, 1997; Sahiner et al., 2000). Finally, the implications of data handling and CAD evaluation with limited datasets have been discussed in detail in several recent publications (Gur, Wagner, & Chan, 2004; Li & Doi, 2006, 2007).

In contrast, the problem of classification using imbalanced data has attracted less attention. It has been mainly addressed in the literature on machine learning (Barnard & Botha, 1993; Chawla, Bowyer, Hall, & Kegelmeyer, 2002; Elazmeh, Japkowicz, & Matwin, 2006; Japkowicz, 2000; Japkowicz & Stephen, 2002; Maloof, 2003; Weiss & Provost, 2001; Zhou & Liu, 2006). These studies were mostly performed using real life problems, where the effects of particular properties of the training data cannot be easily determined (Chawla et al., 2002; Maloof, 2003; Weiss & Provost, 2001). A study oriented on isolating the effect of some data properties was presented in Japkowicz (2000) and Japkowicz and Stephen (2002). However, the investigators did not include the impact of the number of features in the dataset, correlation among features, and the effect of random sampling from the population. In another study (Weiss & Provost, 2001), the authors evaluated the effect of the extent of data imbalance on classifier performance. However, their study was restricted to the classical C4.5 classifier. Moreover, in general machine learning applications, classification performance is often measured using accuracy as the figure of merit (FOM). Unfortunately, accuracy is not a suitable FOM for medical decision support systems where diagnostic sensitivity and specificity are more clinically relevant and better accepted by the physicians.

Although some CAD researchers have dealt with class imbalance within their own application domain (Boroczky, Zhao, & Lee, 2006), to the best of our knowledge, no systematic evaluation of its effect has been reported from this perspective. The purpose of this investigation is to extend the previously reported studies by providing a more comprehensive evaluation

of the effect of class imbalance in the training dataset for the performance of neural network classifiers in medical diagnosis. This effect is studied in the presence of the following other, commonly occurring limitations in clinical data:

- limited training sample size,
- large number of features, and
- correlation among extracted features.

In this study, we also compare two common class imbalance compensation methods (i) oversampling and (ii) undersampling. Since the study specifically targets CAD applications, two distinct neural network training methods are also investigated. The first method is the traditional backpropagation (BP) and the second one is particle swarm optimization (PSO) with clinically relevant performance criteria. This additional factor will allow us to assess whether the neural network training algorithm has any impact on the conclusions.

The article is organized as follows. Section 2 provides a brief description of clinically relevant FOMs for assessing the performance of classifiers for binary classification problems. Section 3 describes the training algorithms employed in this study. Section 4 provides description of the study design and data. Results and discussion follow in Sections 5 and 6, respectively.

2. A clinically relevant FOM: Receiver Operator Characteristic (ROC) Curve

Traditionally, accuracy has been used to evaluate classifier performance. It is defined as the total number of misclassified examples divided by the total number of available examples for a given operating point of a classifier. For instance, in a 2-class classification problem with two predefined classes (e.g., positive diagnosis, negative diagnosis) the classified test cases are divided into four categories:

- true positives (TP) — correctly classified positive cases,
- true negatives (TN) — correctly classified negative cases,
- false positives (FP) — incorrectly classified negative cases, and
- false negatives (FN) — incorrectly classified positive cases.

Therefore, accuracy is

$$\frac{TP + TN}{TP + TN + FP + FN} \quad (1)$$

This evaluation criterion is of limited use in clinical applications for many reasons. First, accuracy varies dramatically depending on class prevalence and it can be very misleading in clinical applications where the most important class is typically underrepresented. For example, if the prevalence of cancer is 5% in the test dataset (a typical clinical scenario), a classifier that detects 100% of cancer-free cases and 0% of cancer cases achieves a seemingly high 95% accuracy. From a clinical perspective, however, this is an unacceptable performance since all cancer patients are misdiagnosed and thus left untreated. Second, in medical decision making, different types of misclassifications have different costs. For example, in breast cancer diagnosis, a false positive decision translates into an unnecessary

breast biopsy, associated with both emotional and financial cost. False negative decision, however, means a missed cancer which in turn can be deadly. Such differences are not taken into account by accuracy. Finally, accuracy depends on the classifier's operating threshold. Since many classification systems (such as neural networks) provide a decision variable of multiple possible values, choosing the optimal decision threshold can be challenging. It also makes impossible direct comparisons among CAD systems that are designed to operate with different decision thresholds.

To account for these issues, Receiver Operator Characteristic (ROC) analysis is commonly used in the clinical CAD community (Obuchowski, 2003). ROC curve describes the relation between two indices: true positive fraction (TPF) and false positive fraction (FPF), defined as follows:

$$\text{TPF} = \frac{\text{TP}}{\text{TP} + \text{FN}} \quad (2)$$

$$\text{FPF} = \frac{\text{FP}}{\text{TN} + \text{FP}}. \quad (3)$$

A conventional ROC curve plots TPF (or sensitivity) vs. FPF (or $[1 - \text{specificity}]$) for every possible decision threshold imposed on the decision variable. By providing such a complete picture, ROC curves are often used to select the optimal decision threshold by maximizing any pre-selected measure of clinical efficacy (e.g., accuracy, average benefit, etc.).

In CAD studies, the most commonly used FOM is the area under the ROC curve (AUC). The AUC index for useful classifiers is constrained between 0.5 (representing chance behavior) and 1.0 (representing perfect classification performance). CAD classifiers are typically designed to maximize the ROC area index. In cancer screening applications, it is expected that the CAD classifier achieves sufficiently high sensitivity (e.g., 90%). Accordingly, researchers have proposed the partial AUC index ($_p\text{AUC}$, where p indicates the lowest acceptable sensitivity level) as a more meaningful FOM (Jiang, Metz, & Nishikawa, 1996). Detailed description of ROC analysis and its utilization for CAD evaluation can be found in Bradley (1997), Jiang et al. (1996) and Metz, Herman, and Shen (1998).

3. Training algorithms for neural networks

Training feedforward neural networks is an optimization problem of finding the set of network parameters (weights) that provide the best classification performance. Traditionally, the backpropagation method (Rumelhart, Hinton, & Williams, 1986) is used to train neural network classifiers. This method is a variation of the gradient descent method to find the minimum of an error function in the weight space. The error measure is typically mean squared error (MSE). Although there is a correlation between MSE and the classification performance of the classifier, there is no simple relation between them. In fact, it is possible that an MSE improvement (decrease) may cause a decline in classification performance. Moreover, MSE is very sensitive to class imbalances in the data. For example,

if positive training examples are severely underrepresented in the training dataset, an MSE-trained classifier will tend to assign objects to the negative class. As described before, such classification is of no use in the clinical domain.

To overcome this limitation, a particle swarm optimization (PSO) algorithm with clinically relevant objectives is also implemented for training neural network classifiers (Habas, Zurada, Elmaghraby, & Tourassi, 2007). The study compares the results of training classifiers using BP method with MSE as an objective with those using the PSO algorithm.

Particle swarm optimization (Kennedy & Eberhart, 1995) is an iterative optimization algorithm inspired by the observation of collective behavior in animals (e.g. bird flocking and fish schools). In PSO, each candidate solution to the optimization problem of a D -variable function is represented in one particle. Each particle i is described by its position x_i (a D -dimensional vector representing a potential solution to the problem) and its velocity v_i . The algorithm typically starts with a random initialization of the particles. Then, in each iteration, the particles change their position according to their velocity. In each iteration the velocity is updated. Given that p_i is the best position (i.e. one that corresponds to the best value of the objective function) found by an individual i in all the preceding iterations and p_g is the best position found so far by the entire population, the velocity of a particle changes according to the following formula (Clerc & Kennedy, 2002; van den Bergh & Engelbrecht, 2004)

$$v_{id}(t) = wv_{id}(t-1) + \varphi_1(t)(p_{id}(t-1) - x_{id}(t-1)) + \varphi_2(t)(p_{gd}(t-1) - x_{id}(t-1)), \quad d = 1, \dots, D, \quad (4)$$

where $v_{id}(t)$ is the d th component of the velocity vector of a particle i in iteration t (analogous notation is used for x_i , p_i , and p_g), $\varphi_1(t)$ and $\varphi_2(t)$ are random numbers between 0 and c_1 and 0 and c_2 , respectively, and w is an inertia coefficient. c_1 , c_2 and w are parameters of the algorithm deciding on significance of particular factors while adjusting the velocity. Position of the particle i is simply adjusted as

$$x_{id}(t) = x_{id}(t-1) + v_{id}(t), \quad d = 1, \dots, D. \quad (5)$$

The output of the algorithm is the best global position found during all iterations. Even though PSO convergence to a global optimum has not been proven for the general case (some results on the convergence can be found in Clerc and Kennedy (2002)), the algorithm has been shown efficient for many optimization problems including training neural networks (Kennedy & Eberhart, 1995).

From the clinical perspective, the most attractive aspect of PSO-based training is that it can be conducted using clinically relevant evaluation criteria. This means that the objective function for the PSO algorithm can be chosen to be AUC, $_p\text{AUC}$ or other clinically relevant criteria (e.g., specific combinations of desired sensitivity and specificity). In this study, the PSO-based neural network training consists of finding the set of weights that provides the best classification performance in terms of AUC or $_{0.9}\text{AUC}$.

Applying ROC-based evaluation during the classifier training could provide multiple benefits. First, since the final evaluation criterion fits the training criterion, the overall performance of the classifier can be potentially improved. Second, since AUC is basically independent of class prevalence, dataset imbalance is of lower concern when training the neural network with clinically relevant objectives. PSO has been successfully applied in CAD for training classifiers with ROC-based objectives (Habas et al., 2007), but its effectiveness with imbalanced datasets has not yet been evaluated.

4. Study design

The study is designed to assess systematically the impact of imbalanced training data on classifier performance while taking into account other factors such as the size of the training dataset, the number of features available, and the presence of correlation among features. The study was conducted with simulated data and the conclusions were further validated using a clinical dataset for breast cancer diagnosis.

The neural networks used in the study were feedforward neural networks with a single output neuron and one hidden layer consisting of three neurons. A network with three hidden neurons was chosen to keep the network complexity low and to prevent overtraining. Sigmoidal activation functions were used for all neurons. The neural networks were trained using (i) BP with MSE, (ii) PSO with AUC and (iii) PSO with $_{0.9}$ AUC as the objective functions. When applying the BP method, all the networks were trained for 1000 iterations with a learning rate of 0.1. For the PSO training, the following standard algorithm parameters were used: $c_1 = 2$, $c_2 = 2$, $w = 0.9$. The number of particles was set to the number of parameters of each neural network (varying based on the number of features) multiplied by 10. The number of iterations was set to 100. The parameters for this study were chosen empirically to provide good performance while at the same time keeping the time complexity feasible.

To prevent overtraining, the examples available for the development of the network were divided into two sets: a training set and a validation set. The training and validation sets were characterized by the same size and positive class prevalence and both were used to construct a classifier. Although choosing equal-sized training and validation sets is unusual (usually a validation set is substantially smaller), it was necessary due to the class imbalance factor. For instance, given a training set with 100 examples and 1% prevalence of positive examples, choosing a validation set smaller than the training would result in no positive validation examples. The training set was used to calculate the MSE and the gradient for BP and to calculate AUC or $_{0.9}$ AUC for the PSO-based training. During the training process, classifier performance on the validation set was repeatedly evaluated. The network that provided the best performance on the validation set during training was selected at the end of the training process. Such practice was applied to prevent possible overfitting of the network to the training examples.

To obtain an accurate estimation of the network performance, a hold-out technique was applied in which a separate set of testing examples (not used in the training) was used to evaluate the network after the training process. For the BP method, one network was trained and finally tested using both FOMs (AUC and $_{0.9}$ AUC). For the PSO training method, a separate classifier was trained for each FOM separately and each trained network was tested on the final test set according to its corresponding FOM.

To account for the class imbalance, two standard ways of compensation were evaluated, namely oversampling and undersampling. In oversampling, examples from the underrepresented class are copied several times to produce a balanced dataset. This means that the size of the training dataset increases up to two times. Note that in the case of batch BP training, this method is equivalent to the commonly used approach where the changes of weights induced by a particular example are adjusted according to the prevalence of its class in the training set (lower prevalence, higher weight change). Also, note that oversampling has no effect on the PSO training as class prevalence does not affect the ROC-based assessment. In undersampling, examples from the overrepresented class are randomly removed, resulting in a smaller dataset. Although the computational complexity of the training process for this method decreases, the main drawback is that potentially useful examples are discarded from training. In both scenarios, the resulting datasets are characterized by equal class prevalence.

4.1. Experiment 1: Simulated data

In the first experiment, simulated data were generated to evaluate the combined effect of all examined factors on classifier performance. Such an evaluation would not be possible with data coming from a real clinical problem since the proposed experiments require a very large number of examples. Furthermore, in simulated data, important parameters can be strictly controlled which allows assessing their separate as well as combined impact on classifier performance. In Experiment 1 we followed the experimental design similar to the one presented in Sahiner et al. (2000).

The simulated datasets were generated using multivariate normal distributions separately for each class:

$$p(\mathbf{x}) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp \left[-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^t \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right], \quad (6)$$

where $p(\mathbf{x})$ is a probability density function, \mathbf{x} is a M -dimensional vector of features, $\boldsymbol{\mu}$ is a vector of means and Σ is a covariance matrix. Furthermore, it was assumed that the covariance matrices for both classes are equal ($\Sigma = \Sigma_1 = \Sigma_2$). Based on the above assumptions, the best achievable AUC performance is given by the following equation (Sahiner et al., 2000):

$$A_z(\infty) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\sqrt{\Delta(\infty)/2}} e^{-t^2/2} dt, \quad (7)$$

where $\Delta(\infty)$ is the Mahalanobis distance between the two classes:

$$\Delta(\infty) = (\mu_2 - \mu_1)^T \Sigma^{-1} (\mu_2 - \mu_1). \quad (8)$$

To evaluate how different levels of imbalances in training dataset affect performance in the presence or absence of feature correlation and for different number of features, two general cases were considered: (i) uncorrelated features and (ii) correlated features. For each of these cases three data distributions were created with 5, 10 and 20 features. For each of the resulting combinations, the two classes had multivariate Gaussian distributions with unequal means and equal covariance matrices.

4.1.1. Distributions for uncorrelated data

For this scenario, it was assumed that the covariance matrices were the identity matrices ($\Sigma_1 = \Sigma_2 = I$) and the difference of means between the classes for feature i was

$$\Delta\mu(i) = \mu_2(i) - \mu_1(i) = \alpha(M - i + 1), \quad i = 1, \dots, M, \quad (9)$$

where M is the number of features and α is a constant. A similar distribution of $\Delta\mu(i)$ was observed in the clinical data used in this study. The parameter α was selected separately for each number of features to provide $\Delta(\infty) = 3.28$ which corresponds to the ideal observer performance $A_z = 0.9$.

4.1.2. Distributions for correlated data

In this scenario, it was also assumed that the covariance matrices for the two classes are equal ($\Sigma = \Sigma_1 = \Sigma_2$) but they are not identity matrices. For each of the number of features M , a 5×5 matrix A_M was constructed. Then the covariance matrix was generated as a block-diagonal matrix based on A_M , i.e., a matrix that has A_M on its diagonal and matrices containing zeros outside the diagonal. For example, the matrix for 10 features was

$$A_{10} = \begin{bmatrix} 1 & 0.1 & 0.2 & 0.3 & 0.1 \\ 0.1 & 1 & 0.7 & -0.3 & 0.4 \\ 0.2 & 0.7 & 1 & -0.1 & 0.3 \\ 0.3 & -0.3 & -0.1 & 1 & 0.2 \\ 0.1 & 0.4 & 0.3 & 0.2 & 1 \end{bmatrix}. \quad (10)$$

And the corresponding covariance matrix Σ was

$$\Sigma = \begin{bmatrix} A_{10} & 0 \\ 0 & A_{10} \end{bmatrix}. \quad (11)$$

Note that for 5 features $\Sigma = A_5$. The values on the diagonal of A_M were always ones. The values outside the diagonal were selected such that for each number of features the classwise correlations averaged 0.08 with standard deviation of 0.2. The correlations in general varied between -0.3 and 0.8 . These values were selected to reflect the correlation structure observed in the clinical data used in this study.

Mean differences between the two classes were selected using Eq. (9). The parameter α was selected to provide ideal observer performance $A_z = 0.9$ for each number of features.

4.1.3. Other data parameters

The positive class prevalence index c was defined as

$$c = \frac{N_{\text{pos}}}{N_{\text{tot}}}, \quad (12)$$

where N_{pos} is the number of positive examples and N_{tot} is the total number of examples in the training dataset. Six levels of c were used: 0.01, 0.02, 0.05, 0.1, 0.2 and 0.5 where the last one corresponds to the equal prevalence of both classes. Positive class prevalence described the extent of imbalance in the training dataset. Additionally, two sizes of the training dataset were investigated (1000 and 100 examples).

4.1.4. Neural network training and testing

Neural networks were trained for all possible combinations of the described factors. For each combination, 50 training and validation datasets were independently drawn from a given distribution and a separate set of neural networks was trained to account for data variability and random factors inherent in neural network training. For each pair of training and validation datasets, the BP training was conducted three times: (i) with original data, (ii) with oversampled data, and (iii) with undersampled data. The PSO-based training was conducted six times for each one of the following combinations: 3 compensation schemes (oversampling + undersampling + no compensation) \times 2 neural networks (one trained using AUC + one trained using 0.9AUC as the training objective).

For the final evaluation, a separate dataset of 10,000 test examples was created. This set was drawn from the same distribution as the training and validation sets (once for each pair of distributions). Such large testing sample size was used to minimize the uncertainty of the classifier's performance estimation. The testing sets were characterized by equal class prevalence (5000 positive and 5000 negative examples). To compare the results for different scenarios a t-test with no assumption about equal variances was applied.

4.2. Experiment 2: Breast cancer diagnosis data

For further validation, the real life problem of breast cancer diagnosis was also studied. Specifically, the problem was to assess the malignancy status of a breast mass. The diagnosis is made based on clinical and image findings (i.e., features extracted by physicians from mammograms and sonograms and clinical features from the patient's history). The original data used in this experiment consisted of 1005 biopsy-proven masses (370 malignant and 645 benign). Each mass was described by a total number of 45 features. The data used in this experiment is an extended version of the data described in detail in [Jesneck, Lo, and Baker \(2007\)](#). It was collected at Duke University Medical Center according to an IRB-approved protocol.

The original set of 1,005 masses was resampled to obtain training sets that reflected the class imbalance simulated in Experiment 1. Due to limited sample size, only one size of the training dataset was investigated. The training and validation sets consisted of 200 examples each throughout the entire experiment. This number was selected to ensure that

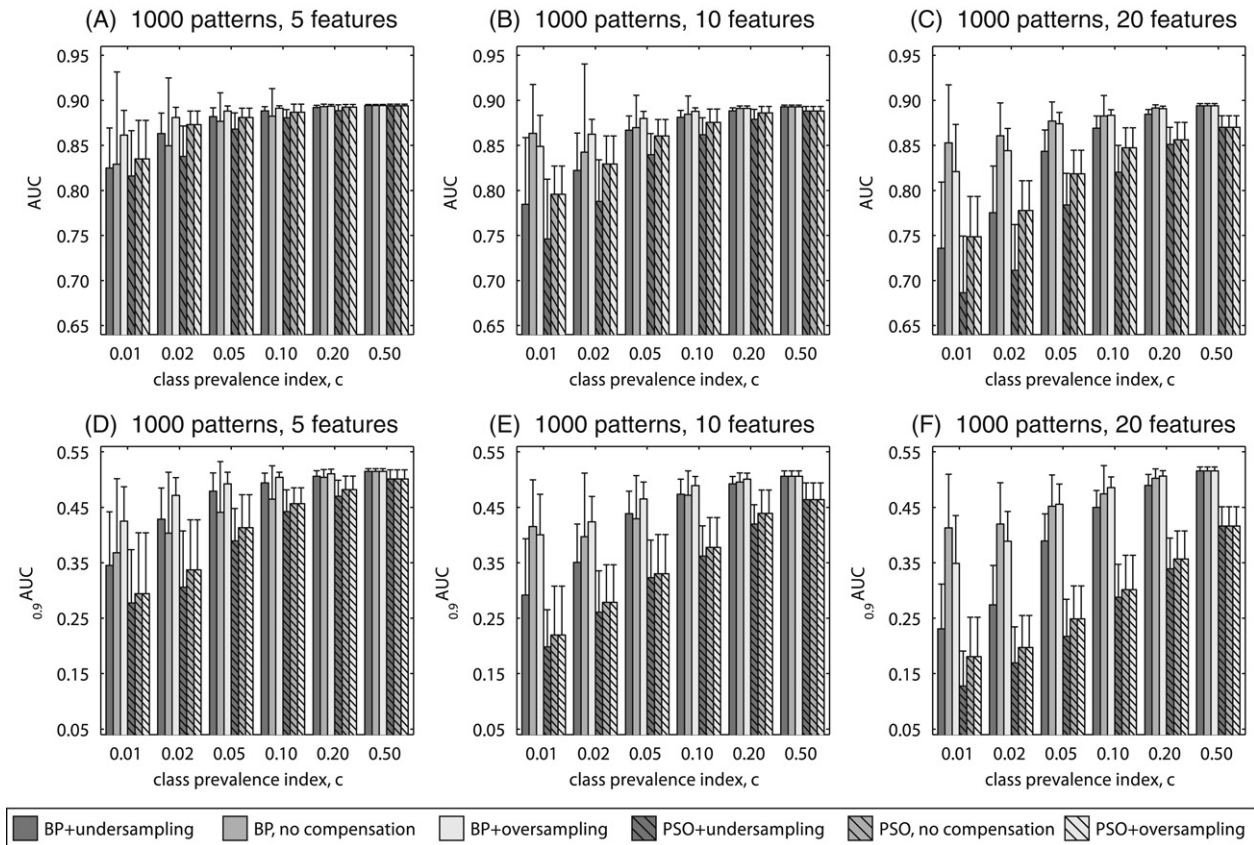


Fig. 1. Simulated data with uncorrelated features: Average testing performance according to AUC and $_{0.9}$ AUC for 1000 training examples.

a sufficient number of cases are excluded for final testing to reduce the estimation variance in testing performance. Actually, 415 examples were excluded for testing. The test set was fairly balanced with 41% cancer prevalence (170 malignant and 245 benign masses). The number of test examples was kept constant so that the variability of the classifier performance will be similar across all studied combinations of parameters. Furthermore, with these examples excluded, there were still enough left to obtain 200 training and 200 validation examples for all class imbalance scenarios considered in this analysis. As in Experiment 1, six values of positive class prevalence were used ranging from 1% to 50%. Three numbers of features were used in this experiment: the original 45 features, 10 features and 5 features. The features were selected using simple forward selection based on the linear discriminant performance with the entire set of 1005 examples. Although it has been clearly shown that feature selection should be done independently of training to avoid an optimistic bias (Sahiner et al., 2000), our study simulates the scenario where diagnostic significance of the particular features is previously known. Studying the impact of class prevalence on feature selection extends beyond the scope of this article. As with the simulated data, for each value of positive class prevalence and number of features, the data was split 50 times to account for the variability introduced by the data split and the stochastic nature of the neural network training.

5. Results

5.1. Experiment 1: Simulated data

The discussion of the study findings is organized around the three main issues: (i) effect of class imbalance on classifier performance, (ii) comparison of neural network training methods and (iii) comparison of data imbalance compensation schemes. The combined effects of data parameters such as number of features and feature correlation are also addressed within the context of the three main issues.

The results of Experiment 1 for uncorrelated features are summarized in Figs. 1 and 2, each showing the neural network average test performance based on the size of the training dataset (1000 and 100 examples, respectively). The error bars show the standard deviations in performance obtained for 50 neural networks. For each figure, there are 6 subplots. The subplots show average values of the two clinical FOMs (top row — AUC, bottom row — $_{0.9}$ AUC) for two different methods of training (bars with no stripes — BP, bars with stripes — PSO) and three class imbalance handling schemes (dark grey — undersampling, medium grey — no compensation, light grey — oversampling). In each row, the three subplots show the results for different number of features: 5 (subplots A and D), 10 (subplots B and E) and 20 (subplots C and F). Selected results for correlated features are shown in Fig. 3 to highlight trends observed as with uncorrelated features.

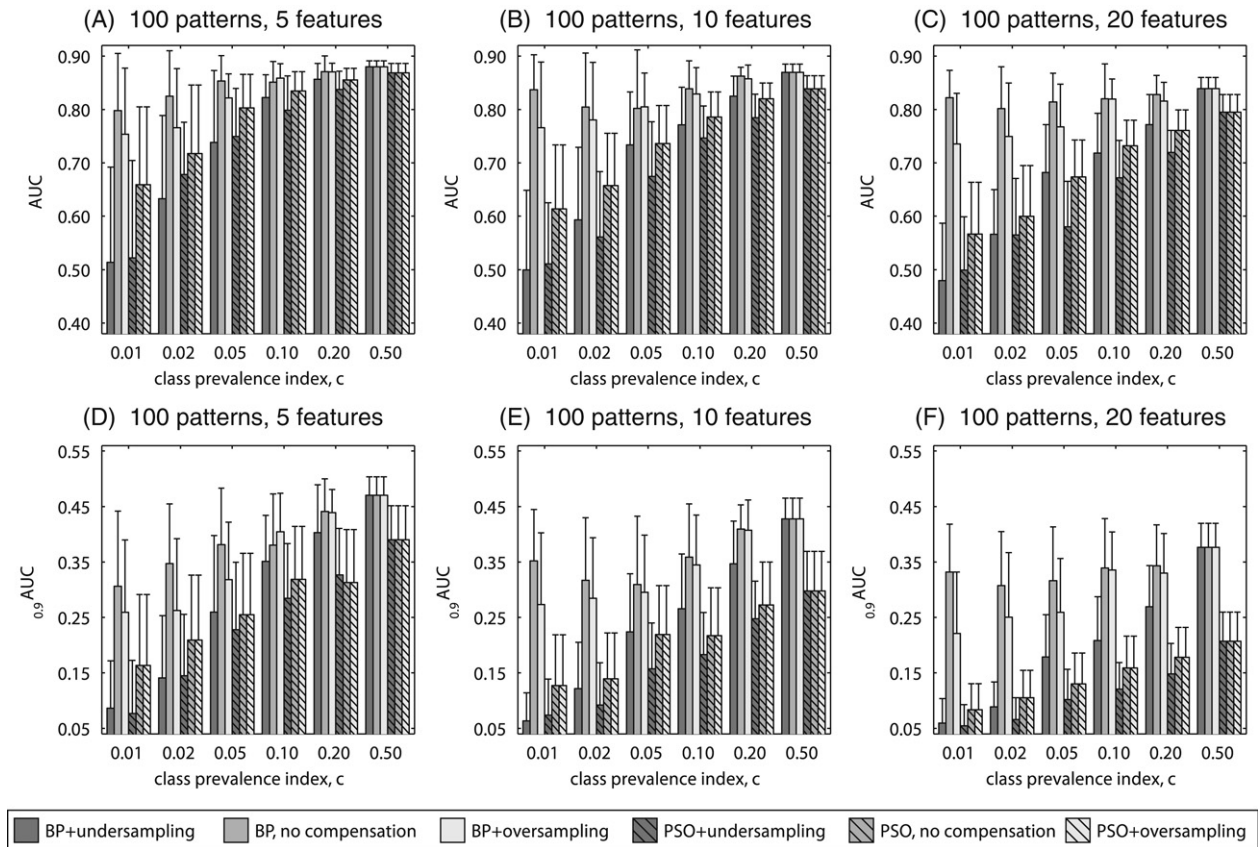


Fig. 2. Simulated data with uncorrelated features: Average testing performance according to AUC and $_{0.9}AUC$ for 100 training examples.

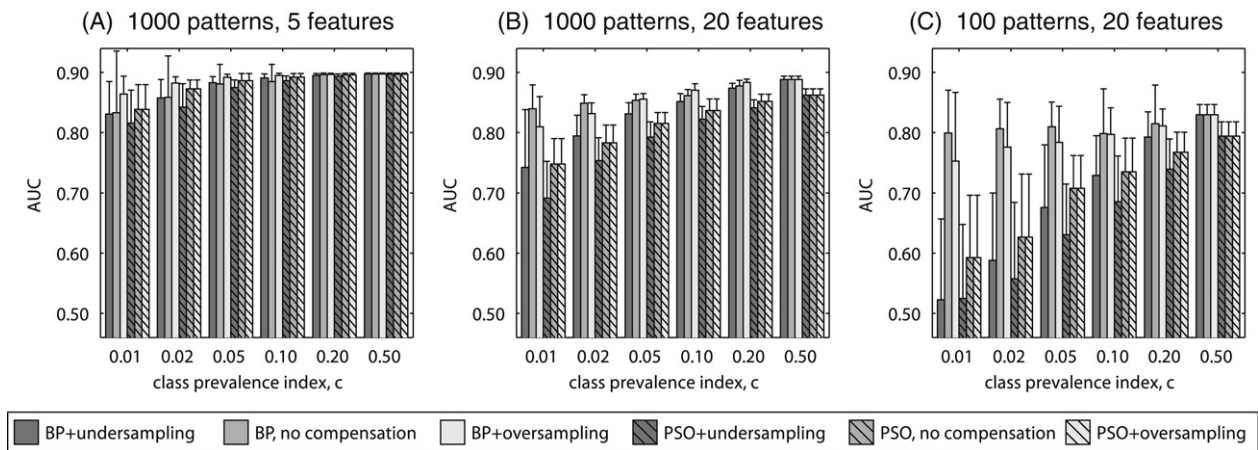


Fig. 3. Simulated data with correlated features: Average testing performance for selected scenarios.

5.1.1. Overall effect of class imbalance

Figs. 1–3 show that, in general, the increasing extent of class imbalance in the training data (i.e., reducing positive class prevalence) has an increasingly detrimental effect on the neural network classifier performance. In fact, for both BP and PSO training and no data imbalance compensation scheme, a statistically significant decline of AUC performance was observed (two-tailed p -value < 0.05) even for small imbalances ($c \leq 0.2$) in vast majority of the cases. Note that the average

performance of the neural networks trained for imbalanced datasets were compared to the performance obtained by the same training method when the data was balanced. This finding was consistent for both 100 and 1000 examples. Class imbalance appears to have even larger effect on the performance measured by $_{0.9}AUC$. For this FOM, a statistically significant decline was observed for class imbalance indices $c \leq 0.2$, for all, except one, analyzed choice of number of features and training sample size. Further, class imbalance appears to result

in larger decline of the average performance when PSO is applied. All these findings were consistent for both correlated and uncorrelated data.

As expected, the decreasing positive class prevalence affected also the variability of the classifier's final performance due to the random sampling of examples from the population and random factors present in the neural network training process. This was expressed in increasing standard deviation of the performance of 50 trained networks when c decreased. For example, for 5 uncorrelated features, 1000 training examples and no correlation, standard deviations of AUC estimates for BP with no compensation increased from 0.001 for $c = 0.5$ to 0.03 for $c = 0.1$ and to 0.1 for $c = 0.01$. With 100 training examples, the standard deviation of AUC increased from 0.01 for $c = 0.5$ to 0.04 for $c = 0.1$ and to 0.11 for $c = 0.01$.

5.1.2. Effect of training method with no compensation

Comparison of the results for the two training methods leads to the following conclusions. For small number of features, PSO and MSE provide similar results for balanced datasets. It can be seen that the average performance for both training methods reaches values close to the population AUC of 0.9. When the number of features increases, average testing performance obtained by both training schemes decreases as well. The decrease in performance, however, is larger for the PSO-based training. Consistently with many previous studies, a low number of training examples has also detrimental effect for the average performance. Again, the drop in performance is larger for PSO training especially for ${}_{0.9}\text{AUC}$. Finally, as stated above the detrimental effect of low class prevalence was also higher for PSO. In conclusion, even though PSO neural network training was shown to be efficient for some tasks, in the scenarios analyzed in this study, BP is a preferable choice in terms of average performance. The results show no clear relation between the training method and performance variance.

5.1.3. Effect of class imbalance compensation scheme

A clear conclusion from the obtained results is that undersampling is not a good choice of compensating for the imbalance in the training data. In fact, in most of the analyzed scenarios, undersampling provided worse performance than both no compensation and oversampling for any type of training and any FOM. This result was consistent for uncorrelated and correlated data.

As mentioned earlier, oversampling has no effect on the PSO training. Comparing oversampling for the BP training and no compensation does not lead to straightforward conclusions. It can be seen in Figs. 1–3 that there are scenarios showing a beneficial effect of the oversampling and cases where oversampling has in fact a detrimental effect on the performance. Some regularities can be observed. Overall, for the examined distributions, oversampling is preferable when the ratio of the number of training examples to number of features is high. For example, it can be seen that oversampling outperforms the no compensation approach for all c for the scenario with 5 features and 1000 training examples. On the other hand, in the case with 20 features and 100 examples,

no compensation should be chosen for all c . This general observation holds for uncorrelated and correlated data and both examined FOMs. Comparison of the oversampling and no compensation approaches in terms of variance shows that generally the method providing a better average performance also results in lower performance variance.

5.2. Experiment 2: Breast cancer diagnosis data

Fig. 4 shows the results of the second experiment based on real clinical data. Most of the general conclusions drawn with experiment 1 hold for the clinical data study as well. Some minor differences, however, must be noted. The number of features has a smaller detrimental effect on the performance of PSO-trained classifiers. The general preference, however of choosing BP-trained classifiers still holds.

In this experiment, a larger impact of class prevalence on the performance is observed. It can be explained by a more complex data distribution and larger number of examples needed to appropriately sample the feature space. Regarding compensation technique for the clinical data, oversampling generally performed better than the no compensation approach. However, there were still some cases where no compensation slightly outperformed oversampling. These results are consistent with the general conclusions drawn when using simulated data.

6. Conclusions

In this study the effect of class imbalance in training data on the performance was evaluated for neural network-based classifiers and the two-class classification problem. The confounding effects of other factors such as training sample size, number of features, and correlation between features were also considered. An extensive experimental study was performed based on simulated dataset and the conclusions were further validated with clinical data for breast cancer diagnosis.

The general conclusions drawn from this study are as follows. First, increasing class imbalance in the training dataset generally has a progressively detrimental effect on the classifier's test performance measured by AUC and ${}_{0.9}\text{AUC}$. This is true for small and moderate size training datasets that contain either uncorrelated or correlated features. In the majority of the analyzed scenarios backpropagation provided better results as PSO training was more susceptible to factors such as class imbalance, small training sample size and large number of features. Again, this finding was true for both correlated and uncorrelated features.

Although undersampling was typically an inferior choice to compensate for class imbalance, there is no clear winner between oversampling and no compensation. The classifier designer should take into account factors such as class distribution, class prevalence, number of features and available training sample size when choosing a compensation strategy for training sets with class imbalances.

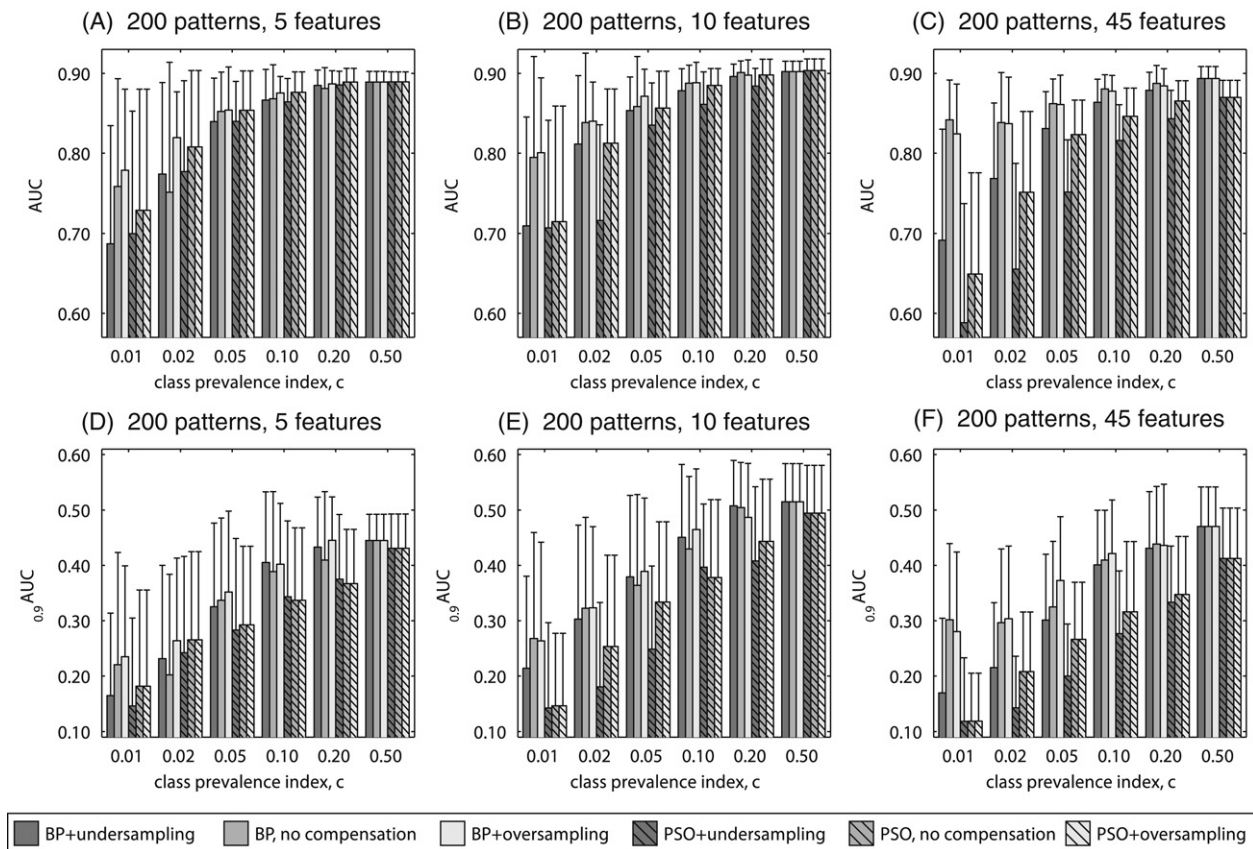


Fig. 4. Results for clinical data. Average testing performance according to AUC and $_{0.9}$ AUC for 200 training examples.

Acknowledgments

This work was supported in part by grants R01-CA-1901911, R01-CA-112437, and R01-CA-95061 from the National Cancer Institute and the University of Louisville Grosscurth Fellowship.

The authors would like to thank the members of the Computational Intelligence Laboratory at the University of Louisville and the members of the Duke Advanced Imaging Laboratory at Duke University, especially Dr. Robert Saunders for helpful discussions.

References

- Aha, D. W., Kibler, D., & Albert, M. K. (1991). Instance-based learning algorithms. *Machine Learning*, 6, 37–66.
- Barnard, E., & Botha, E. C. (1993). Back-propagation uses prior information efficiently. *IEEE Transactions on Neural Networks*, 4, 794–802.
- Beiden, S. V., Maloof, M. A., & Wagner, R. F. (2003). A general model for finite-sample effects in training and testing of competing classifiers. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 23, 1561–1569.
- Boroczky, L., Zhao, L., & Lee, K. P. (2006). Feature subset selection for improving the performance of false positive reduction in lung nodule CAD. *IEEE Transactions Information Technology in Biomedicine*, 10, 504–511.
- Bradley, A. P. (1997). The use of the area under the roc curve in the evaluation of machine learning algorithms. *Pattern Recognition*, 30, 1145–1159.
- Chan, H.-P., Sahiner, B., & Hadjiiski, L. (2004). Sample size and validation issues on the development of CAD systems. In *Proceedings of the 18th international congress and exhibition on computer assisted radiology and surgery*.
- Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). Smote: synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16, 321–357.
- Clerc, M., & Kennedy, J. (2002). The particle swarm – explosion, stability, and convergence in a multidimensional complex space. *IEEE Transactions on Evolutionary Computation*, 6, 58–73.
- Duda, R. O., Hart, P. E., & Stork, D. G. (2000). *Pattern classification*. Wiley-Interscience.
- Elazmeh, W., Japkowicz, N., & Matwin, S. (2006). Evaluating misclassifications in imbalanced data. *Lecture Notes in Computer Science*, 4212, 126–137.
- Fukunaga, K., & Hayes, R. R. (1989). Effect of sample size in classifier design. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 11, 873–885.
- Gur, D., Wagner, R. F., & Chan, H.-P. (2004). On the repeated use of databases for testing incremental improvement of computer-aided detection schemes. *Academic Radiology*, 11, 103–105.
- Habas, P. A., Zurada, J. M., Elmaghraby, A. S., & Tourassi, G. D. (2007). Particle swarm optimization of neural network CAD systems with clinically relevant objectives. In *Proceedings of medical imaging 2007: Computer-aided diagnosis* (pp. 65140M).
- Hamamoto, Y., Uchimura, S., & Tomita, S. (1996). On the behavior of artificial neural network classifiers in high-dimensional spaces. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 18, 571–574.
- Jain, A., & Zongker, D. (1997). Feature selection: evaluation, application, and small sample performance. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 19, 153–158.
- Japkowicz, N. (2000). Learning from imbalanced data sets: a comparison of various strategies. In *AAAI workshop on learning from imbalanced data sets* (pp. 00–05).
- Japkowicz, N., & Stephen, S. (2002). The class imbalance problem: a systematic study. *Intelligent Data Analysis Journal*, 6, 429–450.

- Jesneck, J. L., Lo, J. Y., & Baker, J. A. (2007). Breast mass lesions: computer-aided diagnosis models with mammographic and sonographic descriptors. *Radiology*, 244, 390–398.
- Jiang, Y., Metz, C. E., & Nishikawa, R. M. (1996). A receiver operating characteristic partial area index for highly sensitive diagnostic tests. *Radiology*, 201, 745–750.
- Kawamoto, K., Houlihan, C. A., Balas, E. A., & Lobach, D. F. (2005). Improving clinical practice using clinical decision support systems: a systematic review of trials to identify features critical to success. *British Medical Journal*, 330, 765–772.
- Kennedy, J., & Eberhart, R. (1995). Particle swarm optimization. In *Proceedings of IEEE international conference on neural networks* (pp. 1942–1948).
- Li, Q., & Doi, K. (2006). Reduction of bias and variance for evaluation of computer-aided diagnostic schemes. *Medical Physics*, 33, 868–875.
- Li, Q., & Doi, K. (2007). Comparison of typical evaluation methods for computer-aided diagnostic schemes: Monte Carlo simulation study. *Medical Physics*, 34, 871–876.
- Lisboa, P. J. (2002). A review of evidence of health benefit from artificial neural networks in medical intervention. *Neural Networks*, 15, 11–39.
- Lisboa, P. J., & Taktak, A. F. G. (2006). The use of artificial neural networks in decision support in cancer: a systematic review. *Neural Networks*, 19, 408–415.
- Maloof, M. A. (2003). Learning when data sets are imbalanced and when costs are unequal and unknown. In *Proceedings of workshop on learning from imbalanced data sets*.
- Mazurowski, M. A., Habas, P. A., Zurada, J. M., & Tourassi, G. D. (2007). Impact of low class prevalence on the performance evaluation of neural network based classifiers: Experimental study in the context of computer-assisted medical diagnosis. In *Proceedings of international joint conference on neural networks* (pp. 2005–2009).
- Metz, C. E., Herman, B. A., & Shen, J.-H. (1998). Maximum likelihood estimation of receiver operating characteristic (ROC) curves from continuously-distributed data. *Statistics in Medicine*, 17, 1033–1053.
- Mitchell, T. (1997). *Machine learning*. McGraw Hill.
- Obuchowski, N. A. (2003). Receiver operating characteristic curves and their use in radiology. *Radiology*, 229, 3–8.
- Raudys, S. (1997). On dimensionality, sample size, and classification error of nonparametric linear classification algorithms. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 19, 667–671.
- Raudys, S. J., & Jain, A. K. (1991). Small sample size effects in statistical pattern recognition: recommendations for practitioners. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 13, 252–264.
- Rumelhart, D. E., Hinton, G. E., & Williams, R. J. (1986). Learning internal representations by error propagation. In *Parallel distributed processing: Explorations in the microstructure of cognition, volume 1* (pp. 318–362). MIT Press.
- Sahiner, B., Chan, H. P., Petrick, N., Wagner, R. F., & Hadjiiski, L. (2000). Feature selection and classifier performance in computer-aided diagnosis: the effect of finite sample size. *Medical Physics*, 27, 1509–1522.
- Sampat, M. P., Markey, M. K., & Bovik, A. C. (2005). Computer-aided detection and diagnosis in mammography. In *Handbook of image and video processing* (pp. 1195–1217). Academic Press.
- van den Bergh, F., & Engelbrecht, A. P. (2004). A cooperative approach to particle swarm optimization. *IEEE Transactions on Evolutionary Computation*, 8, 225–239.
- Wagner, R. F., Chan, H.-P., Sahiner, B., & Petrick, N. (1997). Finite-sample effects and resampling plans: applications to linear classifiers in computer-aided diagnosis. In *Proceedings of medical imaging 1997: Image processing* (pp. 467–477).
- Weiss, G. M., & Provost, F. (2001). The effect of class distribution on classifier learning: an empirical study. Technical report, Department of Computer Science, Rutgers University.
- Zhang, G. P. (2000). Neural networks for classification: a survey. *IEEE Transactions on Systems, Man, and Cybernetics – Part C: Applications and Reviews*, 30, 451–462.
- Zhou, Z.-H., & Liu, X.-Y. (2006). Training cost-sensitive neural networks with methods addressing the class imbalance problem. *IEEE Transactions on Knowledge and Data Engineering*, 18, 63–77.